

AWS Glue and Amazon Athena Laboratory Report

Section 3: Laboratory Objectives

General Objective:


Implement a cloud-based data processing solution using AWS Glue, AWS CloudFormation, Amazon S3, and Amazon Athena to analyze weather datasets stored in AWS infrastructure.

Specific Objectives:

- Configure AWS Glue crawlers and databases to process structured weather datasets.
- Deploy cloud infrastructure using AWS CloudFormation templates.
- Execute SQL queries in Amazon Athena to analyze large-scale climate data.

Competencies Developed:

- AWS cloud service integration and administration.
- Infrastructure as Code (IaC) using AWS CloudFormation.
- Data analysis and querying using Amazon Athena.



```
14 # Create a database to contain tables created by the crawler
15 CFNDatabaseWeather:
16   Type: AWS::Glue::Database
17   Properties:
18     CatalogId: !Ref AWS::AccountId
19     DatabaseInput:
20       Name: !Ref CFNDatabaseName
21       Description: "AWS Glue container to hold metadata tables for the weather crawler"
22 #Create a crawler to crawl the weather data on a public S3 bucket
23 CFNCrawlerWeather:
24   Type: AWS::Glue::Crawler
25   Properties:
26     Name: !Ref CFNCrawlerName
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

Section 4: Laboratory Development

4.1 Environment Preparation

Screenshot 1: Initial Configuration

![AWS IAM policy and Glue configuration]

Description:

The initial environment configuration consisted of preparing IAM permissions and validating access to AWS Glue, Athena, and S3 services. IAM policies were reviewed to guarantee the required privileges for database creation, crawler execution, and Athena query execution.

Commands Executed:

Command 1

```
aws configure
```

Command 2

```
aws cloudformation create-stack --stack-name gluecrawler --template-body  
file://template.yaml
```

Command 3

ID de cuenta: 0228...
voclabs/user435464...

IAM > Políticas > Policy-For-Data-Scientists

Denegación explícita (3 de 463 servicios)

Servicio	Nivel de acceso	Recurso
Athena	Limitado: Escribir	Multiple
Glue	Limitado: Administración de permisos, Leer, Escribir	Todos los recursos
S3	Limitado: Administración de permisos, Escribir, Etiquetado	Multiple

Permitir (3 de 463 servicios)

Mostrar los 460 servicios restantes

Servicio	Nivel de acceso	Recurso
Athena	Completo: Etiquetado Limitado: Enumerar, Leer, Escribir	Multiple
Glue	Completo: Enumerar, Etiquetado Limitado: Administración de permisos, Leer, Escribir	Todos los recursos
S3	Limitado: Enumerar, Administración de permisos, Leer, Escribir, Etiquetado	Multiple

Screenshot 2: CloudFormation Stack Deployment

Process Description:

AWS CloudFormation was used to automate the deployment of AWS Glue resources. The YAML template defined the Glue database, crawler configuration, S3 target paths, and schema policies. This approach reduced manual configuration and ensured reproducibility.

Code Evidence:

CFNDatabaseWeather:

Type: AWS::Glue::Database

Properties:

CatalogId: !Ref AWS::AccountId

DatabaseInput:

Name: !Ref CFNDatabaseName

CFNCrawlerWeather:

Type: AWS::Glue::Crawler

Properties:

Name: !Ref CFNCrawlerName

Result Obtained:

The CloudFormation stack was created successfully, generating the AWS Glue database and crawler resources automatically.

4.3 Step 2: AWS Glue Crawler Configuration

Screenshot 3: AWS Glue Crawler Running

![AWS Glue crawler execution]

Process Description:

The AWS Glue crawler was configured to scan the NOAA weather dataset stored in Amazon S3. The crawler automatically inferred the schema structure and generated metadata tables in the Glue Data Catalog.

Applied Configurations:

The screenshot shows the AWS Cloud9 IDE interface. The left sidebar displays the file explorer with a project named 'Cloud9 lr' containing files: 'glue list-crawle', 'gluecrawler.cf.', and 'README.md'. The main editor window is titled 'gluecrawler.cf.yml' and shows a CloudFormation template. The template defines two resources: 'CFNDatabaseWeather' (an AWS::Glue::Database) and 'CFNCrawlerWeather' (an AWS::Glue::Crawler). The crawler is configured with a name, role, and targets pointing to an S3 bucket containing NOAA weather data. The database is configured with a catalog ID and a description. The bottom pane shows the rendered JSON output for the 'Crawler' resource.

```
14 # Create a database to contain tables created by the crawler
15 CFNDatabaseWeather:
16   Type: AWS::Glue::Database
17   Properties:
18     CatalogId: !Ref AWS::AccountId
19     DatabaseInput:
20       Name: !Ref CFNDatabaseName
21       Description: "AWS Glue container to hold metadata table
22 #Create a crawler to crawl the weather data on a public S3 buc
23 CFNCrawlerWeather:
24   Type: AWS::Glue::Crawler
25   Properties:
26     Name: !Ref CFNCrawlerName
```

```
"Crawler": {
  "Name": "cfn-crawler-weather",
  "Role": "gluelab",
  "Targets": {
    "S3Targets": [
      {
        "Path": "s3://noaa-ghcn-pds/csv/by_year/",
        "Exclusions": []
      }
    ],
    "JdbcTargets": [],
    "MongoDBTargets": [],
    "DynamoDBTargets": [],
    "CatalogTargets": [],
    "DeltaTargets": [],
    "IcebergTargets": [],
    "HudiTargets": []
  },
  "DatabaseName": "cfn-database-weather",
  "Description": "AWS Glue crawler to crawl weather data",
  "Classifiers": [],
  "RecrawlPolicy": {
    "RecrawlBehavior": "CRAWL EVERYTHING"
  },
  "SchemaChangePolicy": {
```

Parameter	Value	Justification
Crawler Name	weatherdata	Unique identifier for the crawler
Data Source	NOAA weather dataset in S3	Source of climate records
Recrawl Policy	CRAWL EVERYTHING	Ensures all files are analyzed

Parameter	Value	Justification
Schema Change Policy	Merge	Allows schema evolution

Result Obtained:

The crawler executed correctly and generated the table structure required for Athena queries.

4.4 Step 3: Table Validation in AWS Glue

Screenshot 4: Glue Table Schema

![AWS Glue table schema validation]

Process Description:

After the crawler completed execution, the generated schema was reviewed in AWS Glue. Columns such as station, date, type, observation, and quality flags were detected automatically.

Code Evidence:

```
DESCRIBE by_year;
```

Result Obtained:

The Glue catalog correctly recognized all fields contained in the weather dataset.

4.5 Step 4: Query Execution in Amazon Athena

Screenshot 5: Athena Query Results

![Amazon Athena query results]

Process Description:

Amazon Athena was used to query the weather dataset directly from Amazon S3 using SQL syntax. Queries were executed to analyze temperature and precipitation information across multiple years.

Code Evidence:

```
SELECT date, type, observation
```

FROM by_year

WHERE date/10000 BETWEEN 1950 AND 2015;

Result Obtained:

Athena successfully processed the query and returned weather records from the selected date range.

4.6 Step 5: Climate Data Analysis

Screenshot 6: Average Temperature Results

![[Average weather statistics generated in Athena]]

Process Description:

Advanced SQL queries were executed to calculate average maximum temperatures grouped by year. Athena processed large datasets efficiently without requiring dedicated infrastructure.

Code Evidence:

SELECT year,

AVG(max) AS avg_temperature


FROM weather_data

GROUP BY year

ORDER BY year;

Result Obtained:

The resulting dataset displayed annual average temperatures between 1950 and 1959



.

Section 5: Results and Validation

5.1 Operational Evidence

Final Screenshot: System in Operation

![Complete AWS Glue and Athena workflow operating]

Metrics Obtained:

- Processing time: 29.849 seconds
- Data volume processed: 103.55 GB
- Resources used: AWS Glue, Amazon Athena, Amazon S3, AWS CloudFormation
- Error rate: 0%

Validation Summary:

The implemented infrastructure successfully ingested, cataloged, and queried large-scale weather datasets. Athena queries returned expected results while Glue crawlers generated accurate metadata.

Section 6: Analysis and Reflection


6.1 Challenges Encountered

Challenge 1:

- Description: Incorrect IAM permissions initially prevented Glue services from accessing S3 resources.
- Solution: IAM policies were modified to allow Glue and Athena access to the required buckets and services.
- Learning: Proper IAM configuration is essential for secure and functional AWS environments.

Challenge 2:

- Description: Athena queries initially produced schema mismatches due to metadata inconsistencies.
- Solution: The crawler was re-executed with updated schema change policies.
- Learning: Glue crawlers simplify schema management for evolving datasets.



```
10   Type: String
11   Default: cfn_sample_1-weather
12   # Resources section defines metadata for the Data Catalog
13   Resources:
14   # Create a database to contain tables created by the crawler
15   CFNDatabaseWeather:
16     Type: AWS::Glue::Database
17     Properties:
18       CatalogId: !Ref AWS::AccountId
19       DatabaseInput:
20         Name: !Ref CFNDatabaseName
21         Description: "AWS Glue container to hold metadata table
22   #Create a crawler to crawl the weather data on a public S3 buc
23   CFNCrawlerWeather:
24     Type: AWS::Glue::Crawler
25     Properties:
26       Name: !Ref CFNCrawlerName
27       Role: arn:aws:iam::022854128902:role/gluelab
28       #Classifiers: none, use the default classifier
29       Description: AWS Glue crawler to crawl weather data
30       #Schedule: none, use default run-on-demand
31       DatabaseName: !Ref CFNDatabaseName
32       Targets:
33         S3Targets:
34           # Public S3 bucket with the weather data
35           - Path: "s3://noaa-ghcn-pds/csv/by_year/"
36       TablePrefix: !Ref CFNTablePrefixName
37       SchemaChangePolicy:
38       UpdateBehavior: "UPDATE_IN_DAT
```

Section 7: AWS Costs

Cost Estimation:

Service	Usage	Estimated Cost
EC2	Minimal laboratory use	\$0.50
S3	Weather dataset storage	\$0.20
AWS Glue	Crawler execution	\$1.00
Athena	Query processing (103.55 GB)	\$0.60

Service	Usage	Estimated Cost
CloudFormation	Infrastructure deployment	\$0.00
TOTAL		\$2.30

AWS Cost Explorer Screenshot:

! [Cost estimation evidence from AWS]

Cost Optimization Strategies Applied:

- Reused existing AWS lab resources to reduce infrastructure costs.
- Executed optimized SQL queries to minimize scanned data volume in Athena.

Section 8: Conclusions

Technical Conclusions:

- AWS Glue successfully automated metadata extraction and schema generation.

- Amazon Athena enabled serverless analysis of large datasets using SQL queries.

United States (No)
Account ID: 0228-5...
voclabs/user435464...

[AWS Glue](#) > Crawlers

Crawler successfully starting
×

The following crawler is now starting: "weatherdata"

0
 0
 2
 0
 0

Crawlers

A crawler connects to a data store, progresses through a prioritized list of classifiers to determine the schema for your data, and then creates metadata tables in your data catalog.

Crawlers (1/1)
[Info](#)

Last updated (UTC)
February 25, 2026 at 04:15:49

Action
Run
Create crawler

View and manage all available crawlers.

Name	State	Schedule	Last run	Last run...	Lo
weatherdata	Running	-	-	-	-

Feedback
Privacy
Terms
Cookie preferences

```
4 ) AS SELECT date, type, observation FROM by_year
5 WHERE date/10000 between 1950 and 2015;
```

SQL Ln 3, Col 74



Ejecutar de nuevo ▼

Explicar [↗](#)

Cancelar

Borrar

Crear ▼

☐ Volver a utilizar los
resultados de la consulta
hasta hace 60 minutos [✎](#)

Resultados de la consulta

Estado de la consulta


✔ Completado

Tiempo en cola: 96 ms Tiempo de ejecución: 29.849 sec

Datos analizados: 103.55 GB

La consulta se ha realizado correctamente.

- CloudFormation simplified infrastructure deployment and improved reproducibility.



The screenshot shows the Amazon Athena console interface. At the top, there's a navigation bar with the AWS logo, a search icon, and a user profile. Below this, the breadcrumb navigation shows 'Amazon Athena' and 'Editor de consultas'. The main content area displays 'Resultados (10)' with buttons for 'Copiar' and 'Descargar resultados en formato CSV'. A search bar labeled 'Filas de búsqueda' is present. The results are shown in a table with columns: #, date, type, and observation. The table contains 10 rows of data, all for the date '19630101'. The 'type' column has values 'PRCP' and 'TMAX', and the 'observation' column has values '0', '329', '283', '305', and '430'.

#	date	type	observation
1	19630101	PRCP	0
2	19630101	PRCP	0
3	19630101	PRCP	0
4	19630101	PRCP	0
5	19630101	TMAX	329
6	19630101	TMIN	283
7	19630101	PRCP	0
8	19630101	TAVG	305
9	19630101	PRCP	0
10	19630101	TMAX	430

Practical Applicability:

The laboratory demonstrates how organizations can implement scalable cloud-based analytics pipelines for big data processing without managing physical infrastructure. These techniques are applicable in climate monitoring, business intelligence, and large-scale data engineering projects.

Identified Improvement Areas:

- Area 1: Improve query optimization techniques to reduce Athena processing costs.
- Area 2: Implement automated ETL pipelines with AWS Lambda.

- Area 3: Add dashboards and visualization tools such as Amazon QuickSight.

The screenshot shows the Amazon Athena console interface. At the top, there's a navigation bar with the AWS logo, a home icon, a grid icon, a search icon, a document icon, a bell icon, a question mark icon, a settings icon, and a dropdown menu showing 'Estados Unidos'. On the right, it says 'ID de cuenta: 0228' and 'voclabs/user435464...'. Below the navigation bar, the breadcrumb trail shows 'Amazon Athena > Editor de consultas'. A green status bar indicates the query is 'Completado' (Completed) with a checkmark icon. It also shows 'Tiempo en cola: 110 ms', 'Tiempo de ejecución: 8.623 sec', and 'Datos analizados: 2.26 GB'. Below this, there's a section for 'Resultados (66)' with buttons for 'Copiar' (Copy) and 'Descargar resultados en formato CSV' (Download results in CSV format). A search bar labeled 'Filas de búsqueda' is present. The results are displayed in a table with columns '#', 'Year', and 'Max'. The table shows 10 rows of data, with years ranging from 1950 to 1959 and corresponding 'Max' values.

#	Year	Max
1	1950	16.734313818557784
2	1951	16.806621817314802
3	1952	17.280452854305018
4	1953	17.918841157082877
5	1954	17.478422108186066
6	1955	17.104085210486158
7	1956	17.154974111756808
8	1957	17.404536896943593
9	1958	17.301316467760206
10	1959	17.379689889526748

Section 9: References


[1] Amazon Web Services, "AWS Glue Documentation," AWS Docs. [Online]. Available: <https://docs.aws.amazon.com/glue/> [Accessed: 18-May-2026].

[2] Amazon Web Services, "Amazon Athena Documentation," AWS Docs. [Online]. Available: <https://docs.aws.amazon.com/athena/> [Accessed: 18-May-2026].

[3] Amazon Web Services, "AWS CloudFormation Documentation," AWS Docs. [Online]. Available: <https://docs.aws.amazon.com/cloudformation/> [Accessed: 18-May-2026].

[4] Apache Software Foundation, "Apache Hive Documentation," Version 4.0. [Online]. Available: <https://hive.apache.org/> [Accessed: 18-May-2026].

[5] National Oceanic and Atmospheric Administration, "Global Historical Climatology Network Daily," NOAA. [Online]. Available: <https://www.ncei.noaa.gov/> [Accessed: 18-May-2026].



The screenshot shows the Amazon Athena console interface. At the top, there's a navigation bar with the AWS logo, a search icon, and a user profile section showing 'ID de cuenta: 0228' and 'voclabs/user435464'. Below this, the breadcrumb navigation indicates 'Amazon Athena > Editor de consultas'. The main content area displays 'Resultados (10)' with two buttons: 'Copiar' and 'Descargar resultados en formato CSV'. A search bar labeled 'Filas de búsqueda' is present. The results are shown in a table with columns: '#', 'date', 'observation', and 'type'. The table contains 10 rows of data, all with the date '19630101' and the type 'TMAX'. The 'observation' values range from 106 to 194. The table has a light gray header and alternating row colors for readability.

#	date	observation	type
1	19630101	106	TMAX
2	19630101	78	TMAX
3	19630101	89	TMAX
4	19630101	156	TMAX
5	19630101	178	TMAX
6	19630101	189	TMAX
7	19630101	189	TMAX
8	19630101	183	TMAX
9	19630101	167	TMAX
10	19630101	194	TMAX

Section 10: Code Repository

GitHub Repository:

- URL: [https://github.com/\[username\]/aws-glue-athena-weather-lab](https://github.com/[username]/aws-glue-athena-weather-lab)
- Branch: main
- Commit Hash: latest-commit

Repository Structure:

project/

└─ README.md

└─ src/

| └─ cloudformation/

| └─ queries/

| └─ config/

└─ config/

| └─ glue-config/

└─ docs/

| └─ screenshots/

└─ tests/

| └─ integration/

└─ data/

| └─ sample/

└─ requirements.txt

└─ .gitignore

Execution Instructions:

Clone repository

git clone [repository-url]

Install dependencies

```
pip install -r requirements.txt
```

```
# Configure environment
```

```
aws configure
```

```
# Run deployment
```

```
aws cloudformation deploy --template-file template.yaml --stack-name gluecrawler
```

Section 11: Appendices

Appendix A: Complete Logs

2026-02-24 22:08:16 - CloudFormation stack creation started

2026-02-24 22:10:42 - AWS Glue crawler created successfully

2026-02-25 04:15:49 - AWS Glue crawler execution started

2026-02-25 04:18:12 - Athena query completed successfully

Appendix B: Complete Code

Resources:

CFNDatabaseWeather:

Type: AWS::Glue::Database

CFNCrawlerWeather:

Type: AWS::Glue::Crawler

Appendix C: Additional Screenshots

- IAM policy configuration
- Athena query execution
- Glue crawler running state
- Glue schema validation

Appendix D: Configuration Files

RecrawlPolicy:

RecrawlBehavior: CRAWL_EVERYTHING

SchemaChangePolicy:

UpdateBehavior: MERGE_NEW_COLUMNS